

Contextual Advertising
Apache Hive and Amazon EMR

Storing Logs on Amazon S3

```
s3://elasticmapreduce/samples/hive-ads/tables/impressions/  
dt=2009-04-13-08-05/ec2-12-64-12-12.amazon.com-2009-04-13-08-05.log
```

Launching a Development Job Flow

```
ssh -i ~/my-keypair-private-key.pem hadoop@ec2-67-202-12-120.compute-1.amazonaws.com
```

Running a Hive Session on the Master Node

```
hive \  
-d SAMPLE=s3://elasticmapreduce/samples/hive-ads \  
-d DAY=2009-04-13 -d HOUR=08 \  
-d NEXT_DAY=2009-04-13 -d NEXT_HOUR=09 \  
-d OUTPUT=s3://test-data-services
```

```
ADD JAR ${SAMPLE}/libs/jsonserde.jar ;
```

Declaring Tables in Amazon S3

```
CREATE EXTERNAL TABLE impressions (  
  requestBeginTime string, adId string, impressionId string, referrer string,  
  userAgent string, userCookie string, ip string  
)  
PARTITIONED BY (dt string)  
ROW FORMAT  
  serde 'org.apache.hive.hcatalog.data.JsonSerDe'  
  with serdeproperties ( 'paths'='requestBeginTime, adId, impressionId, referrer, userAgent, userCookie, ip' )  
LOCATION '${SAMPLE}/tables/impressions' ;
```

```
ALTER TABLE impressions ADD PARTITION (dt='2009-04-13-08-05') ;  
MSCK REPAIR TABLE impressions
```

```
CREATE EXTERNAL TABLE clicks (  
  impressionId string  
)  
PARTITIONED BY (dt string)  
ROW FORMAT  
  SERDE 'org.apache.hive.hcatalog.data.JsonSerDe'  
  WITH SERDEPROPERTIES ( 'paths'='impressionId' )  
LOCATION '${SAMPLE}/tables/clicks' ;
```

```
MSCK REPAIR TABLE clicks
```

Combining the Clicks and Impressions Tables

```
CREATE EXTERNAL TABLE joined_impressions (
  requestBeginTime string, adId string, impressionId string, referrer string,
  userAgent string, userCookie string, ip string, clicked Boolean
)
PARTITIONED BY (day string, hour string)
STORED AS SEQUENCEFILE
LOCATION '${OUTPUT}/joined_impressions'
;

CREATE TABLE tmp_impressions (
  requestBeginTime string, adId string, impressionId string, referrer string,
  userAgent string, userCookie string, ip string
)
STORED AS SEQUENCEFILE;

INSERT OVERWRITE TABLE tmp_impressions
SELECT
  from_unixtime(cast((cast(i.requestBeginTime as bigint) / 1000) as int)) requestBeginTime,
  i.adId, i.impressionId, i.referrer, i.userAgent, i.userCookie, i.ip
FROM
  impressions i
WHERE
  i.dt >= '${DAY}-${HOUR}-00' and i.dt < '${NEXT_DAY}-${NEXT_HOUR}-00'
;

CREATE TABLE tmp_clicks (
  impressionId string
) STORED AS SEQUENCEFILE;

INSERT OVERWRITE TABLE tmp_clicks
SELECT
  impressionId
FROM
  clicks c
WHERE
  c.dt >= '${DAY}-${HOUR}-00' AND c.dt < '${NEXT_DAY}-${NEXT_HOUR}-20'
;

INSERT OVERWRITE TABLE joined_impressions PARTITION (day='${DAY}', hour='${HOUR}')
SELECT
  i.requestBeginTime, i.adId, i.impressionId, i.referrer, i.userAgent, i.userCookie,
  i.ip, (c.impressionId is not null) clicked
FROM
  tmp_impressions i LEFT OUTER JOIN tmp_clicks c ON i.impressionId = c.impressionId
;
```

Running in Script Mode

s3://elasticmapreduce/samples/hive-ads/libs/join-clicks-to-impressions.q

```

$ SAMPLE=s3://elasticmapreduce/samples/hive-ads
$ OUTPUT=s3://mybucket/samples/output
$ ./elastic-mapreduce --create --name "Join Clicks" \
  --hive-script --arg $SAMPLE/libs/join-clicks-to-impressions.q \
  --args -d,SAMPLE=$SAMPLE \
  --args -d,DAY=2009-04-13,-d,HOUR=08 \
  --args -d,NEXT_DAY=2009-04-13,-d,NEXT_HOUR=09 \
  --args -d,INPUT=$SAMPLE/tables \
  --args -d,OUTPUT=$OUTPUT \
  --args -d,LIB=$SAMPLE/libs

```

Contextual Advertising Model

to estimate the probability of a click given the context.
 $P[\text{click}|\text{context}]$

One heuristic for doing this is the following formula.
 $\text{product}_{\{f \text{ in context}\}} \Pr[\text{click}|f=\text{true}]$

This heuristic multiplies the probability of a click for each feature that is true in the advertising context. If we take the negative log of this formula, we get the following formula.
 $-\sum_{\{f \text{ in context}\}} \log \left(\frac{\text{count}[\text{click}, f=\text{true}]}{\text{count}[f=\text{true}]} \right)$

Declaring External Tables in the Interactive Job Flow

```

hadoop@domU-12-31-39-07-D2-14:~$ hive \
  -d SAMPLE=s3://elasticmapreduce/samples/hive-ads

```

```

CREATE EXTERNAL TABLE IF NOT EXISTS MODEL_joined_impressions (
  request_begin_time string, ad_id string, impression_id string,
  page string, user_agent string, user_cookie string, ip_address string,
  clicked boolean
)
PARTITIONED BY (day STRING, hour STRING)
STORED AS SEQUENCEFILE
LOCATION '${SAMPLE}/tables/joined_impressions';

```

```
MSCK REPAIR TABLE MODEL_joined_impressions;
```

```
SHOW PARTITIONS MODEL_joined_impressions;
```

Producing the Feature Matrix

transformation on our impression data to produce Boolean features

User Agent

#_____

convert the user agent string into a sequence of keywords is to use a python script. As we'll see shortly, we can call this script directly from within a Hive statement.

```
#!/usr/bin/python
```

```
import sys
import re
```

```
for line in sys.stdin:
    user_agent, ad, clicked = line.strip().split('\t')
    components = re.split('[:;,\(\)]', user_agent)
    for component in components:
        if len(component) != 0:
            print '\t'.join([component, ad, clicked])
```

```
# call this script from within a Hive, we issue a MAP statement.
```

```
MAP
MODEL_joined_impressions.user_agent, MODEL_joined_impressions.ad_id,
MODEL_joined_impressions.clicked
USING
'${SAMPLE}/libs/split_user_agent.py' AS
feature, ad_id, clicked
FROM
MODEL_joined_impressions
LIMIT 10;
```

<code_snippet>

Query ID = hadoop_20210316030721_9f7629e0-cf49-4676-94ea-027fc85b66a1

Total jobs = 1

Launching Job 1 out of 1

Status: Running (Executing on YARN cluster with App id application_1615776600752_0017)

```
-----
      VERTICES   MODE      STATUS TOTAL COMPLETED RUNNING PENDING FAILED KILLED
-----
Map 1 ..... container  SUCCEEDED   12      12      0      0      0      0
-----
VERTICES: 01/01 [=====>>] 100% ELAPSED TIME: 21.76 s
-----
```

OK

```
Mozilla u375QMdkb3BaP0Kv4QaQTFoxg587nV true
5.0      u375QMdkb3BaP0Kv4QaQTFoxg587nV true
Windows  u375QMdkb3BaP0Kv4QaQTFoxg587nV true
U        u375QMdkb3BaP0Kv4QaQTFoxg587nV true
Windows  u375QMdkb3BaP0Kv4QaQTFoxg587nV true
NT       u375QMdkb3BaP0Kv4QaQTFoxg587nV true
5.1      u375QMdkb3BaP0Kv4QaQTFoxg587nV true
en-US    u375QMdkb3BaP0Kv4QaQTFoxg587nV true
rv:1.9.0.10 u375QMdkb3BaP0Kv4QaQTFoxg587nV true
Gecko    u375QMdkb3BaP0Kv4QaQTFoxg587nV true
Time taken: 23.352 seconds, Fetched: 10 row(s)
```

</code_snippet>

The columns user_agent, ad_id, and clicks from the joined_impressions table are input to the script and the result is a table with the columns feature, ad_id, and clicked.

The output of the statement is displayed on the console so we limit the number of lines output to ten. We can see from the output that the keywords contain spaces and are not lower cased. To normalize the output we

apply the user defined functions trim and lower and we prefix each keyword by 'ua:' so these features can be mixed with other features.

```
SELECT concat('ua:', trim(lower(temp.feature))) as feature, temp.ad_id, temp.clicked
FROM (
  MAP MODEL_joined_impressions.user_agent, MODEL_joined_impressions.ad_id,
MODEL_joined_impressions.clicked
  USING '${SAMPLE}/libs/split_user_agent.py' as feature, ad_id, clicked
  FROM MODEL_joined_impressions
) temp
LIMIT 10;
```

IP Address

```
SELECT
  concat('ip:', regexp_extract(ip_address, '^([0-9]{1,3}\.){0-9}[0-9]{1,3}).*', 1)) AS
  feature, ad_id, clicked
FROM
  MODEL_joined_impressions
LIMIT 10;
```

OK

ip:45.206	AAv8arUW6Uw8HsXfssxearjVRbIOU9	false
ip:43.201	S15U6hxbUmrFNdJvQLi9KtQDRlwkv	false
ip:43.189	mgH3jRI4TKjU0pwNCHTu062sMUKMi7	false
ip:32.32	le4T9koKQPnHjWKixqrR6ofwDSotrW	false
ip:21.45	LVTkiC1ThKHde4JjmkfrGqnigMVUeM	false
ip:34.192	d1M9EmKxHfLsOb5xG94X6vWVfXtGMm	false
ip:61.87	GOr35CvAg0dGE4QcQE8ChJMPRMpc0o	false
ip:20.198	VgsaJ62M6D7r0apuCSpmalQ1d66qQg	true
ip:66.205	rRbJTRwfru8lpTMsDteTOIhU984Vji	false
ip:40.175	d1M9EmKxHfLsOb5xG94X6vWVfXtGMm	false

Time taken: 1.155 seconds, Fetched: 10 row(s)

URL

#_____

To extract a feature from the URL of the page on which the advertisement displays, we make the URLs all lowercase and add "page:" to the beginning.

```
SELECT concat('page:', lower(page)) as feature, ad_id, clicked
FROM MODEL_joined_impressions
LIMIT 10;
```

OK

page:yimg.com	AAv8arUW6Uw8HsXfssxearjVRbIOU9	false
page:schuelervz.net	S15U6hxbUmrFNdJvQLi9KtQDRlwkv	false
page:21cn.com	mgH3jRI4TKjU0pwNCHTu062sMUKMi7	false
page:ibibo.com	le4T9koKQPnHjWKixqrR6ofwDSotrW	false
page:barnesandnoble.com	LVTkiC1ThKHde4JjmkfrGqnigMVUeM	false
page:paipai.com	d1M9EmKxHfLsOb5xG94X6vWVfXtGMm	false
page:startimes2.com	GOr35CvAg0dGE4QcQE8ChJMPRMpc0o	false
page:indiatimes.com	VgsaJ62M6D7r0apuCSpmalQ1d66qQg	true
page:gougou.com	rRbJTRwfru8lpTMsDteTOIhU984Vji	false
page:gougou.com	d1M9EmKxHfLsOb5xG94X6vWVfXtGMm	false

Time taken: 0.127 seconds, Fetched: 10 row(s)

Combining the Features

Now that we've written queries to normalize each of the feature types let's combine them into one table. We can do this using Hive's UNION operator. Keep in mind that all sub queries in the union must have the same number of columns that have the same, exact names.

```
SELECT *
FROM (
  SELECT concat('ua:', trim(lower(ua.feature))) as feature, ua.ad_id, ua.clicked
  FROM (
    MAP MODEL_joined_impressions.user_agent, MODEL_joined_impressions.ad_id,
MODEL_joined_impressions.clicked
    USING '${SAMPLE}/libs/split_user_agent.py' as (feature STRING, ad_id STRING, clicked BOOLEAN)
    FROM MODEL_joined_impressions
  ) ua
)

UNION ALL

SELECT concat('ip:', regexp_extract(ip_address, '^[0-9]{1,3}\.[0-9]{1,3}.*', 1)) as feature, ad_id, clicked
FROM MODEL_joined_impressions

UNION ALL

SELECT concat('page:', lower(page)) as feature, ad_id, clicked
FROM MODEL_joined_impressions
) temp
limit 50;
```

Note that we had to modify the user agent query slightly. Passing data through a mapper strips the columns of their types and returns them as strings. To merge with the other tables, we need to define clicked as a Boolean.

Index Table

```
CREATE TABLE feature_index (
  feature STRING,
  ad_id STRING,
  clicked_percent DOUBLE )
STORED AS SEQUENCEFILE;

INSERT OVERWRITE TABLE feature_index
SELECT
  temp.feature,
  temp.ad_id,
  sum(if(temp.clicked, 1, 0)) / cast(count(1) as DOUBLE) as clicked_percent
FROM (
  SELECT concat('ua:', trim(lower(ua.feature))) as feature, ua.ad_id, ua.clicked
  FROM (
    MAP MODEL_joined_impressions.user_agent, MODEL_joined_impressions.ad_id,
MODEL_joined_impressions.clicked
    USING '${SAMPLE}/libs/split_user_agent.py' as (feature STRING, ad_id STRING, clicked BOOLEAN)
    FROM MODEL_joined_impressions
  ) ua
)
```

UNION ALL

```
SELECT concat('ip:', regexp_extract(ip_address, '^([0-9]{1,3}\.[0-9]{1,3}).*'), 1)) as feature, ad_id, clicked
FROM MODEL_joined_impressions
```

UNION ALL

```
SELECT concat('page:', lower(page)) as feature, ad_id, clicked
FROM MODEL_joined_impressions
) temp
GROUP BY temp.feature, temp.ad_id;
```

cast(count(clicked = 'true') as DOUBLE)

sum(if(clicked = 'true', 1, 0))

Applying the Heuristic

```
SELECT
  ad_id, -sum(log(if(0.0001 > clicked_percent, 0.0001, clicked_percent))) AS value
FROM
  feature_index
WHERE
  feature = 'ua:safari' OR feature = 'ua:chrome'
GROUP BY
  ad_id
ORDER BY
  value ASC
LIMIT 100
;
```

OK

9DGhVCJPfCkoXA1jJ3UtHGwfU0whlL	0.9808292530117262
k60lJn5sr3JxT9tOgCRA29jgow7CX	1.0986122886681098
IcJq5iHUCXPgMVPfIGfxFC3XhbiRJj	1.2321436812926323
heG6lG11EQwGgIQStm2dv0M6qe3saP	1.252762968495368
RVtOSibJxch3OJoEELE3RMIUC2jVKc	1.2580400255962119
mAvGKX705nFv35rennhwEqL6BGJXQM	1.349926716949016
eEPpmatr12b3IwfGwbGfo2fPDUpIm	1.4226620052907655
7XkCp8ndNPQaH88GTwsReE8qNqlfff	1.4350845252893225
XshiIrb6d7S8IEd25bPIq1trJNauuK	1.4350845252893225
3uk7HODLcmubKbQu7uD HoupTrHv21l	1.4469189829363254
UXEIPrcKVoePINjvhg7BeRj2utWNjc	1.4816045409242156
Qax0xaEx1u44DeJP0OwAPRKQ41QUs9	1.4816045409242156
73oEqx3Mkk88apIpLwp02IgroItcHt	1.5260563034950494
ubKkdWnDqUNjaaH5Mh2tLfHGeO4gkg	1.5686159179138452
4pGQfqJWhTcgNFCmGIqaJjEcipI49n	1.6582280766035324
fe02VgK4oFJp4hcgHSmxthJa9xM3HP	1.6582280766035324
IdW1dwPIBh7WdKCxE3Wjljk06J6vR	1.7047480922384253
7ORhspStIoX59Mh3Io2FLw8NbBLcEA	1.7272209480904839
wuanC2Q3uOFjF2vle1UsDSAO4IaCgv	1.7346010553881064
35o4b04xEPh5I4ts1p4jIAsQkGSR9i	1.7491998548092589
6Puw5UL9LwWdfb01192Jq2mCwfITUB	1.749199854809259
OSk5qN7upGVHjsC8VmTae0W6jkTLiW	1.776011112259916

3Efxu1QEJPi0ES1oPKpUT31wKxgs1C	1.791759469228055
MliHGihLh6V0Ci4oMifpORQXu5pU6m	1.791759469228055
IXkhO5fdhb7c3NlkC9kHP4Sv4uA4aG	1.791759469228055
aJXiRii3hsQ8FrHqWogKRinnixTvhX	1.791759469228055
ihXF7HBq62LK7suwDFguW4qG0TvoGU	1.791759469228055
3LUsLrdCj4sRSDEUNslg7MwXTkaLMx	1.791759469228055
v6GTMOrrSSTseBP50bT8tNHUxLADagj	1.791759469228055
9eO9paWEuuT04S06HTqP8eGpAw3197	1.8137383759468304
FiNO6TnQpg4n03fvv61whuDshJjQwq	1.8325814637483102
13QnG1wqVSj64LJo5RpJuaX5grRVhN	1.8325814637483102
lIXISS5vjcaBpawl1vQV3dfGwI5UxG	1.8458266904983307
8Lcc413EqE2BNbC2kx1aLRJfKlQeKV	1.8458266904983307
I5FhtPH1XepqgkPBBGWhVIne6BpgjW	1.845826690498331
jVKbCxOWq4WEQFVC3hoMffWcho3KwF	1.845826690498331
HXXvt8irJJfQJlJ2HUU9lJfohtnQBG	1.845826690498331
r5JHuEtIOHWonBfcrrVfeHoGHtpS9R	1.845826690498331
lV3T61466Su7d3e722AqnWpvuKNmgh	1.8523840910444898
IKLmMmR04bUvArVJwX1UPvCX9Pi6jS	1.856297990365626
n25ash1v0BG5sX0kFmWdvdIU3UQ8St	1.8718021769015913
mVvHohs0l5ejIraS16mPs652mGQ11O	1.8971199848858813
oCeUObOGV9RV9K9HnVT8ChgCDfaAAAt	1.9095425048844383
lJwD7MewucGb8joT1altmLsehCJkDH	1.916922612182061
dIDkwPFeRxE6dKSUMPOndMAaUajKQT	1.929909807708872
l8WaV2TvtGdURckeMAGEcRacdlu48	1.9315214116032138
sSP4gEKCUhEruOxS8WcOoIKRq6FMBm	1.9459101490553135
2s4geni54iw4QaulC1U5FRf8xH2ePJ	1.9459101490553135
JRN4RBXsC4QPdEDE881uPhAOJk4VpR	1.9459101490553135
9gl50aO3ufeBagOfnm3tW8Ob8bO0kX	1.9459101490553135
49BQTWdeNW7f9XqSxkTNRBf2FS7Wph	1.9459101490553135
KghFbdpcOEed6aKaiEelqqPSpuqXNl	1.9503644994046936
dVS09w0SFNPFq6vXI7BBcxnOTeii5W	1.992430164690206
95JUFwJtrPqffPNVUBCOKN9dNTCh3x	1.9924301646902063
5Ngeh0duRrigoMxUhr8ltoxBG2FDsk	2.0149030205422647
nUCBo2CTdvmkSTQWHBQ80NUpXeJ1Dv	2.0149030205422647
k828HBVrEXUIV6nPR3DhKsxR1dMpk	2.0149030205422647
6tgVI9eJDBPPtwLh3b7CubpC1MnleV	2.0281482472922856
04x4hNC8jQEoP490vrM3Lvr9h61CMX	2.0281482472922856
ssRQ75fpPbslu9IdVPCIE1qNRRon8N	2.0281482472922856
HsodC44GXjGWDfWSAw1Tb7EKruNU2j	2.03688192726104
vo3l6c560k9MliqjsbRXvaFfNiGBkn	2.03688192726104
IuB9hK22dL4cGgcrvEFxSmSECXjCLS	2.043073897508961
opAmlcxjUjclxJvXaedH0VMCKbhAbJ	2.063693184711697
u63dGD1UTBA1RSkCLEkHNf9vgrQ8op	2.0794415416798357
lOt010LorFnmxEETqTCecCpeT3dW1t	2.0794415416798357
PCN4IL8Sl7aemdlJ1dNh8VmAdr6x1	2.0794415416798357
lW30ngNcqL6mLiMRRDsF28OsTtQxTQ	2.0794415416798357
GoAv3IXUTekhIAQOGSCcU2OWBWANND	2.0794415416798357
tlT19bqp99iJXSdruT9LqVSFVvsLvQ	2.0794415416798357
SXToreP7mS4mjcrrw1uCj0HViutjqmC	2.0794415416798357
w46pDADVMEEITLk53x6BR11AHJM5XU	2.0794415416798357
kUmGcbJddc4rBVTITTMqxSIREolkvv	2.0794415416798357
lPseeb3R8E4hwpGag2IPJkpW0u2DS4	2.0794415416798357
Gajk8ugfpiUcDCDN6lnat3RK1Lnlka	2.079441541679836
U5w3KaFRkK99WFLb4XRgfmUQUhlsUw	2.11021320034659
dQWJ2S94CSBxmMO9fqJxMdjlwdnEJG	2.11021320034659
l7xboC6Hu76p5D4Jo2rvHPS4gsPMTH	2.11021320034659
E3N74naULTD2G1PlpiqxopBwXPnEPM	2.120263536200091

ohq16LTNxpw36VRJr4N0ttKcGE4HdB	2.120263536200091
FMdP6don03udLbQAQiPnNo2abhGFeM	2.128231705849268
Vk2doIBUMEKU3vVrbO3k9uDEhvsISL	2.1400661634962708
x9MDjm0BRRf4fSa9skom4QhXcqiDls	2.1400661634962708
8neBI4519203PudofJkaGWKjw9o9T3	2.1400661634962708
92xq1wEUNphQe8dEDITpJ2laRSO5K6	2.1400661634962708
9902DpKFAIQgbaV4KG090Bj9NTXjB1	2.1400661634962708
2ILT8sUXwrHA3BxINf9N1po8o2216d	2.1400661634962708
dvLmOqws1pIJ7IvqwSOCv04QNiHftd	2.1400661634962708
RoMqfrFg6a7Grj18EU3UqmbJ0RJ0go	2.1400661634962708
xmFwSDLCquMnaC2m2CwT8waM1uRSmo	2.1400661634962708
hSNF8n5dglDQfKSRfGqRHDWeWg5luV	2.1546649629174235
8INx50MexumO6DkJFESTcGs3QQW3pS	2.159484249353372
6u7oQe0K0E4oKvjBSHkbktprCggWl	2.174751721484161
EFuNH2tJwpeONN1XAjsDxvaCIdchuI	2.174751721484161
dNbNjB6SgjAWffh5sVkX0OE4aGvU2B	2.182298927119544
TVmW8JfBgOxGEmLpJh5UWdApNurLU	2.197224577336219
5R6sgaxLLW9KaXEMEBecQOruo4BchD	2.1972245773362196
0gvWD0vDouqfVsARqot7jMjK3VXHUL	2.1972245773362196
9cJwSRdGn4UrtmhqIIAfbvS72cDBDI	2.1972245773362196
VLt4V0Lu2sBuxJIMFbxPIT33Iq4K6I	2.1972245773362196

Time taken: 15.632 seconds, Fetched: 100 row(s)

The result is advertisements ordered by a heuristic estimate of the chance of a click. Observing predominance of advertisements for Apple products.

Summary

Developed a job flow to process impression and click logs uploaded to S3 by web server machines.

The result of this job flow is a table in Amazon S3 that is used to develop and test a model for contextual advertising.

The Hive statements collected are used within a job flow to generate a model file.

Uploaded the file to Amazon S3 and thus making it available to adserver machines to serve ads contextually.