

PRASHANT SHISHODIA

• prashantsh0404@gmail.com • linkedin.com/in/prashant-shishodia • github.com/prashantsh-0404 • 503-544-9245 • Portland, OR

EDUCATION

Texas A&M University, Mays Business School

College Station, Texas

Master of Science in Management Information Systems

May 2019

National Institute of Technology, Allahabad

Allahabad, India

Bachelor of Technology in Mechanical Engineering

May 2014

TECHNICAL SKILLS

AWS Analytics: Redshift, EMR, SageMaker, Kinesis, Athena

Programming Languages: Python, SQL, Java, Scala

ETL Processing: Glue, Hadoop, Spark, Kafka, S3, Spectrum

Relational Databases: Amazon RDS, MySQL, PostgreSQL

NoSQL databases: DynamoDB, Neptune, Elasticsearch, DMS

Data Visualization: QuickSight, Kibana, CloudWatch

Compute: EC2, ELB, Lambda, Kubernetes, Containers

Security: IAM, SSO, Cognito, VPC, CloudTrail

WORK EXPERIENCE

Amazon Web Services, Inc.

Portland, USA

Cloud Engineer, Big Data Analytics

July 2019 – present

- Resolved more than 600 cases of production/quality environment issues in cloud infrastructure for AWS customers across Analytics services with customer rating of 5/5; this resulted in cost savings on customer end & consistent improvement in customer satisfaction
- Took ownership of several technical case escalations to assist AWS Service Teams and Account managers to meet immediate requirements of customers like administering cluster with 300+ TBs of data, reducing query execution time, etc.
- Replicated customer-end workload scenarios (e.g., automate ETL pipelines, debug code, eliminate bottlenecks, implement query optimization, etc.) to help customers in building & maintaining AWS architecture acc. to best practices resulting in AWS revenue growth

➤ PROJECT 1: Contextual Advertising using Apache Hive and Amazon Elastic MapReduce (EMR)

- Developed a job flow in EMR to process advertisement impression and click logs generated by EC2 servers and stored in S3
- Created EMR-Hive job to process S3 data, develop and test a model; model file uploaded to S3 & made accessible to Ad Servers
- Using Hive Queries, advertisements are ordered by a heuristic estimate of the chance of a click for contextual advertising

➤ PROJECT 2: Unified serverless streaming ETL architecture with Kinesis Data Analytics (KDA), Amazon RDS, DMS & DynamoDB

- Event data sent to Kinesis data streams using data models in Amazon RDS for MySQL & AWS Database Migration Service
- Processed 1000K order items enriched with order event data & product reference data in KDA app; output sent to DynamoDB
- Observed average latency of 900ms from the time of event ingestion into kinesis pipeline to data records persisted in DynamoDB

➤ PROJECT 3: Organize Amazon Redshift-Based ETL pipelines using serverless AWS Step Functions and Glue Python Shells jobs

- State machine launches Glue Python job tasks which move customer reviews dataset from S3 to Redshift tables via Spectrum layer
- Jobs execute aggregation queries to group reviews by product & ordered by votes; results UNLOAD to S3 for further processing
- Tasks contain error handling switch which, if caught, triggers SNS topic which forward error notifications to subscribers

➤ PROJECT 4: Develop Data Warehouse Architecture and Reports for Retail Chain using SQL Server Management Studio

- Created ETL workflows to process a data load of approx. 20GB and designed data warehouse architecture from scratch
- Developed reports and dashboards to answer business questions in detail; suggested pricing strategies to increase revenue

Larsen and Toubro, Electrical & Automation

Mumbai, India

Business Intelligence Engineer, IT Business Solutions

July 2014 - May 2017

- Designed analytical dashboards & visualizations to calculate *Overall Equipment Efficiency (OEE)* of 12 production machines in IOT project
- Product Owner of Plant Maintenance mobile application involving documentation of business requirements, development and end-to-end implementation; trained users to use mobile app decreasing *Mean Time to Repair (MTR)* of 12 production machines by 20%
- Build predictive and prescriptive analytics models to automate and optimize inventory process metrics; developed inventory forecasting strategies increasing inventory turnover by 6%

LEADERSHIP SKILLS

- Mentored team's 4 new-hires in AWS Big Data & Analytics services using technical expertise and Amazon's Leadership Principles
- Prepared article on Kinesis Streams architecture & got it published under AWS public website's Training documentation